

Sushant Poudel

AI Security & Agent Safety · Applied AI Engineering

Bhaktapur, Nepal | sushant.poudel2028@gmail.com
github.com/sushant-me | linkedin.com/in/sushant-poudel2028 | sushantpoudel2028.com.np | HackingHub: nofear976

WHY THIS PAGE LOOKS LIKE THIS

A CV asks you to take its claims on trust. This one does not. Every entry below names the public source it came from, and all of them are re-checked against those sources once a week by a script in github.com/sushant-me/reputation, which fails the moment one stops being true. Claims that cannot be checked from outside — a signed experience letter, a conference acceptance email — are marked **on request** rather than mixed in with the verified ones.

VERIFIED, WITH THE SOURCE BESIDE IT

Rank #1, HackingHub Q3 2026 global leaderboard — 116 flags, 11,860 XP, 2 silver + 1 bronze award (next account: 97 flags)	api.hackinghub.io/leaderboard · profile
Security patch authored, merged, then generalised into google/go-github — my change stopped a release-asset upload being sent to an off-host URL; the maintainer replaced it the next day with the broader rule "credentials are sent only to configured origins", which is what master implements today and which carries the commit "Address feedback from sushant-me"	#4556 (merged 2026-09-17) · #4564
Five memory-safety issues in google/s2geometry — NULL dereference on the documented traversal path, out-of-bounds read, 16 GiB allocation from a 117-byte input, and a 2.4 GiB allocation reachable from a 28-byte input . Each filed with a standalone reproducer.	issues #674 , #676 , #677 , #678 , #679
Two hardening pull requests under review in <code>google/s2geometry</code> , split into independent changes at the maintainer's request; a third was withdrawn after AddressSanitizer showed it did not fix the crash it claimed to	#681 · #682
MCP tool-shadowing fixes open against three of Google's agent frameworks — framework-reserved tool names that a server-supplied tool can occupy, bypassing the duplicate-name guard	adk-go #1606 · adk-java #1515 · adk-python #7145
An OSS-Fuzz harness that was testing nothing — <code>google/libphonenumber</code> 's as-you-type formatter asserted 0.00% line coverage of the code it targeted because it never generated valid input; the corrected harness reaches 93%	PR #4079
Two IEEE papers accepted — <i>Edge-Native Semantic Firewall for Autonomous LLM Agents</i> (accepted; camera-ready in progress; to be presented at NCIT) and <i>PREBAS: Preemptive Bandwidth Scaling for WebRTC in LEO Satellite Networks</i> (2026 IEEE RTC, Chicago)	paper , code , raw model outputs · PREBAS acceptance on request
Two years as an AI engineer — Atmos SoftTech Pvt. Ltd., Birgunj, Nepal (March 2024 – March 2026): production AI/ML systems, full-stack web delivery, and autonomous agent pipelines including offline-first inference	signed experience letter on request (Md. Samsad Ali, CEO)
Open-source tooling — <code>agentbound</code> (tool-boundary bug detector), <code>mcp-nameguard</code> (MCP tool-name collisions), <code>trajectorycheck</code> (agent trajectory evaluation)	agentbound · mcp-nameguard · trajectorycheck

- **The semantic firewall (600 scenarios per condition, Phi-3-mini 3.8B, 4-bit, 4.2 GiB VRAM budget, zero cloud inference).** Constraining model output to JSON *without* a reasoning field made the evaluator less safe, not more: 46.2% unsafe accepts against 17.2% unconstrained. Mandating a Chain-of-Thought field order cut false accepts of irreversible hard-denial commands from 71 to 6 and reached 66.3% decision accuracy — and still approved 6 of 208 hard denials, which the paper reports rather than rounds away.
- **Measurement discipline as a method.** `beyond-attention` implements a selective state-space model from scratch and publishes the results that do not flatter it — a scalar register still ties the agent on the arithmetic families, and the state-space model is the weaker of the two past its training length.

Check any of this yourself in one command.

```
git clone https://github.com/sushant-me/reputation && cd reputation && python3  
verify_evidence.py
```

It prints one line per claim and exits non-zero if a published claim no longer holds.

Computer Engineering, Nepal Engineering College (final year) · two-time Hult Prize 1st Runner-Up (2024, 2025) · Aspire Leaders Program Global Finalist · certifications: IBM AI Engineering, IBM DevOps, Google Cybersecurity, Google UX Design, DeepLearning.AI Deep Learning, Stanford Machine Learning.

Statuses are stated as their sources report them: an accepted paper with the camera-ready in progress is not called published, and an open pull request is not called merged.